
DAWN-SRW: A Computational Model of Learned State Transitions

Seungho Choi¹

Abstract

Dense language models learn effective state transitions, but the operations that produce those transitions and the conditions governing their use remain largely implicit. This motivates a broad question: *Can learned language-model computation be organized as an explicit state-conditioned operational system?* DAWN-SRW studies a narrower architectural hypothesis: transition-producing computation can be factorized into identifiable learned operations and explicit state-conditioned control that determines how those operations contribute to the subsequent state transition. Here, identifiable means referable as a stable architectural object and available for measurement or intervention, not necessarily human-interpretable. The evaluated v4172 model is one concrete realization. It uses persistent rank-1 read-write (RW) banks, learned operation-address control, continuous-valued applicability coefficients with exact support, and normalized additive composition to construct Q/K/V attention interfaces and direct RST updates; operator visibility and RW execution remain global and dense. At 393.8M parameters and 39.999B tokens, DAWN-SRW retains performance against a matched dense Transformer: C4 cross-entropies are 2.9402 versus 2.9359, and six-task zero-shot means are 0.4629 versus 0.4597. A descriptive 1.272B evaluation reaches 2.8359 cross-entropy and a 0.4858 zero-shot mean. The model exhibits exact input-conditioned applicability support, and a discovery-frozen 4,096-site Q/K-centered operator circuit passes suppression, four matched-control families, and exact restoration on independent validation and held-out indirect-object-identification inputs. Fail-closed results delimit where localization did not succeed. Together, these results provide a constructive test of the hypothesis under the evaluated conditions, without establishing human-readable semantics, general editability, benchmark superi-

ority, or hardware efficiency.

1. Introduction

Modern language models achieve strong performance through repeated dense transformations, but those transformations do not expose the constituent operations or the conditions under which a state uses them. The long-term motivation is to make learned computation available as explicit computational objects that can be identified, analyzed, and causally intervened on. Such an interface could eventually support semantic interpretation, reliable editing, or more programmable learned systems. Those outcomes are motivations, not claims established by this paper.

We begin from a broad question:

Can learned language-model computation be organized as an explicit state-conditioned operational system?

One decomposition of that question is: *How should a language model organize, identify, and apply the operations relevant to its current representation state?*

At a high level, conventional programming makes computation explicit through states, operations, and rules governing their execution. Operational accounts describe execution through transitions between configurations (Plotkin, 1981):

current state $\longrightarrow$ conditions / applicable operations
 $\longrightarrow$ execution $\longrightarrow$ next state.

Such explicit structure supplies objects through which computation can be referred to and analyzed. We use this programming perspective only as an organizational reference, not to claim that language models implement symbolic programs or that DAWN-SRW defines a formal operational semantics.

Language models likewise evolve computational states by repeatedly transforming representations across depth. Let X_ℓ denote the sequence representation at layer ℓ ; at a coarse level, a block implements

$$X_\ell \longrightarrow X_{\ell+1}.$$

¹Independent Researcher. Correspondence to: Seungho Choi <madst0614@gmail.com>.

We likewise use this state-evolution view only as an organizing observation, not as a contribution or a claim about an intrinsic language-model semantics.

In conventional Transformers, the mechanisms producing these state changes remain largely implicit in dense parameterized transformations. Attention and feed-forward layers expose module boundaries, but they do not natively expose a referable set of operations together with explicit state-conditioned control over their application.

We therefore study a narrower architectural hypothesis:

Can transition-producing language-model computation be factorized into identifiable learned operations and explicit state-conditioned control that determines how those operations contribute to the subsequent state transition?

This hypothesis separates what computation is available from when and how it is applied. “Identifiable” means that an operation is referable as a stable architectural object whose contribution can be measured or intervened on; it does not specify how broadly or how long that identity must persist, and does not imply a human-readable concept, a discrete instruction, or a uniquely correct decomposition. The paper provides a constructive architectural test, not evidence that this factorization is intrinsic to all language models.

The evaluated DAWN-SRW v4172 architecture instantiates one point in this broader design space. It chooses a rank-1 read-write transformation as its RW primitive and normalized additive composition as its application semantics. For a token state x ,

$$\mathcal{O}_i(x) = \langle x, r_i \rangle w_i, \quad (1)$$

$$F(x) = \frac{1}{Z(x)} \sum_i c_i(x) \mathcal{O}_i(x), \quad c_i(x) \in [0, 1]. \quad (2)$$

The read direction r_i determines what is extracted from the current representation, the write direction w_i determines the direction of the primitive’s output contribution, and $c_i(x)$ determines whether and to what degree that primitive applies. Thus, this realization parameterizes transition-producing computation through explicit conditional composition of RW primitives.

The control abstraction is also explicit:

$$x \mapsto C_a(x) = (c_{a,1}(x), \dots, c_{a,N_a}(x)). \quad (3)$$

These components instantiate the narrower hypothesis: the representation supplies the current state, $C_a(x)$ controls primitive application, and the composed result supplies computation used in the next update. Continuous-valued

coefficients with exact support and normalized addition are v4172 design choices, not requirements of the architectural hypothesis.

More specifically, v4172 uses persistent type-specific QK, V, and RST RW banks across depth; a state-derived query and threshold τ ; read/write-derived primitive addresses in a learned spherical geometry; exact applicability support with continuous-valued coefficients; normalized additive composition; and deployment across Q, K, V, and the direct residual-state transition (RST). Each successive state recomputes applicability over the corresponding bank via

$$x \mapsto (q_a(x), \tau_a(x)) \mapsto c_{a,i}(x),$$

using primitive addresses $k_{a,i}$. The operation-address space implements control rather than RW execution. The current model evaluates applicability with global operator visibility and applies the RW banks densely, even when exact applicability support is a strict subset of the pool.

This realization builds on the read-write view of dense feed-forward computation, where first-layer directions detect input patterns and second-layer directions contribute outputs (Geva et al., 2021). That analysis reveals functional directions inside an existing dense layer. DAWN-SRW instead turns RW factorization into an architectural constraint: reusable RW units are executable RW primitives whose state-dependent contributions are directly controlled and observable.

The same formulation produces the query, key, and value states used by causal self-attention and the direct update applied to the post-attention residual state. Q and K produce separate outputs from one shared QK pool; V and the post-attention residual transition (RST) use separate pools. Operator pools and applicability projections are reused across depth, while normalization and attention output projections remain layer-specific.

The empirical study follows a fixed claim hierarchy. First, viability is tested by comparing a 393.8M parameter DAWN-SRW checkpoint with a parameter- and token-matched dense Transformer on C4 and six zero-shot tasks, with a descriptive 1.272B evaluation. Second, conditional organization is tested by measuring exact input-conditioned applicability support. Third, causal accessibility is tested by localizing a discovery-only indirect object identification (IOI) circuit. The circuit is frozen before confirmatory data are opened, then tested with suppression, matched controls, and exact restoration on independent validation and fully held-out test inputs. Fourth, three fail-closed boundaries are reported: RAVEL localization, ARC causal validation, and 1.3B IOI localization. Efficiency is not a headline result: the current implementation scores and applies full operator pools densely.

Our contributions are:

1. **An explicit architectural factorization and one realization.** We formulate transition-producing computation in terms of identifiable learned operations and state-conditioned control that determines how those operations contribute to the subsequent state transition. We instantiate that factorization in v4172 through persistent rank-1 RW banks, learned spherical operation-address control, exact applicability support with continuous-valued coefficients, and normalized composition across Q/K/V/RST.
2. **A constructive empirical test with bounded evidence.** We test viability at 400M scale, report a descriptive 1.3B evaluation, characterize conditional organization through applicability support, and test causal accessibility through a discovery-frozen IOI circuit confirmed by suppression, matched controls, and exact restoration. Fail-closed results bound where the same localization program did not succeed.

2. Related Work

2.1. State-dependent and dynamically parameterized transformations

State-dependent parameterizations have long precedents in nonlinear control and neural sequence models. State-dependent Riccati equation techniques use state-dependent coefficient representations in nonlinear control design (Cloutier, 1997); we cite this as conceptual lineage, not as a mathematical equivalence to DAWN-SRW. Multiplicative RNNs allow the current input to determine the hidden-state transition matrix (Sutskever et al., 2011), input-switched affine networks use input-dependent recurrent weights (Forster et al., 2017), and HyperNetworks generate weights for another network (Ha et al., 2017). These works address different settings—feedback control, recurrent dynamics, and generated neural parameters—but collectively establish that a transformation may depend on the current state or input. State evolution, state dependence, and dynamically generated transformations are therefore not contributions of DAWN-SRW. The narrower hypothesis studied here is that transition-producing language-model computation can be factorized into identifiable learned operations and explicit state-conditioned control that determines how those operations contribute to the subsequent state transition. Persistent RW primitives, address-based control, and normalized additive composition are choices of the evaluated realization rather than the hypothesis itself.

2.2. Conditional computation and fine-grained experts

Sparsely-gated mixtures of experts activate a routed subset of expert networks (Shazeer et al., 2017), and Switch Transformers scale this pattern with simple routing (Fedus et al.,

2022). PEER applies product-key retrieval to more than a million tiny experts (He, 2024); OmniMoE extends this fine-grained line with vector-level Atomic Experts, a Cartesian Product Router, and execution scheduling (Shi et al., 2026). Thus state-dependent routing, fine-grained experts, large expert pools, and sparse retrieval all precede DAWN-SRW. These systems primarily develop routing, capacity, and efficient sparse execution for expertized layers. The paper instead studies explicit state-conditioned operator organization as the base parameterization of transition-producing computation, and exposes the resulting operator contributions for causal intervention. Its current implementation still performs a dense address scan and dense, chunked RW application rather than indexed sparse execution.

2.3. Shared computation across depth

Parameter sharing across depth and layer-agnostic expert reuse also have direct precedents. Universal Transformers recur over shared layer parameters (Dehghani et al., 2019), while MoEUT develops a shared-layer Transformer with expertized feed-forward and attention computation (Csordás et al., 2024a). More recent MoE designs move expert capacity beyond layer-local ownership: MoUE reuses a universal layer-agnostic expert pool across layers (Chen et al., 2026), and UniPool exposes one global pool through independent per-layer routers (Huang et al., 2026). The former line studies recurrence or shared-layer computation, while the latter relaxes layer-local expert ownership. DAWN-SRW therefore does not claim depth sharing or a global pool as new. It persistently reuses fine-grained, type-specific QK, V, and RST RW banks; each depth’s current representation recomputes applicability over the same bank, and the shared operation capacity constructs Q/K/V attention interfaces or a direct residual transition.

2.4. Conditional attention and unified expertized transformations

Conditional computation has also been applied beyond Transformer feed-forward layers. SwitchHead routes among attention-side experts and reduces the number of attention matrices evaluated (Csordás et al., 2024b). UMoE reformulates attention to expose an FFN-like expert structure, then studies a shared expert set across attention and FFN computation (Yang et al., 2025). They establish that expertized transformations and parameter sharing need not be confined to the FFN. These works rule out positioning attention-side experts or conditional attention itself as a DAWN-SRW contribution. DAWN-SRW instead applies one state-conditioned RW operator-composition abstraction consistently to the learned Q, K, and V interfaces and the direct post-attention residual-state transition, with persistent type-specific RW banks reused across depth.

2.5. Read-write primitives and memory views

Transformer feed-forward layers admit a key-value-memory interpretation in which read-side directions detect input patterns and write-side directions contribute outputs (Geva et al., 2021). Product Key Memory factorizes retrieval keys (Lample et al., 2019), and scalable memory layers extend large content-addressable trainable memories (Berges et al., 2024). Fast Weight Programmers provide an algebraic and programming precedent based on additive outer-product updates (Schlag et al., 2021). MoRAM explicitly treats rank-1 adapters as associative-memory experts and uses their intrinsic keys for self-activation (Lu et al., 2026). Together these lines establish separate precedents for read/write interpretation, scalable retrieval, outer-product algebra, and rank-1 memory experts. RW, rank-1, and memory views are therefore not new individually. DAWN-SRW uses RW factorization as an executable architectural primitive rather than only a post-hoc interpretation: the architecture directly exposes each primitive’s state-dependent applicability, executed response, and contribution to the composed computation.

2.6. Neural programming and programming interpretations

Neural Programmer-Interpreters represent and execute programs from supervised traces (Reed & de Freitas, 2016), while Neural Programmer differentially selects operations and data (Neelakantan et al., 2016). Fast Weight Programmers connect outer-product computation to elementary instructions (Schlag et al., 2021). DAWN-SRW does not induce a symbolic program, define a formal operational semantics, or require individual operators to be semantically interpretable. Here “program-like” refers only to explicit learned operations, state-conditioned applicability, and their controlled application in producing state transitions. Plotkin’s operational perspective is an organizational reference rather than a formal semantics for the model.

2.7. Mechanistic analysis and causal intervention

Mechanistic interpretability studies behavior through localized components and causal interventions. IOI work isolates an attention-head circuit and evaluates faithfulness, completeness, and minimality (Wang et al., 2023). AtP* provides a scalable approximation to activation patching for localizing behavior to model components (Kramár et al., 2024). RAVEL shows why neuron-level representations may remain entangled and evaluates methods against multiple causal criteria (Huang et al., 2024). This literature distinguishes a candidate localization from evidence that an intervention causally controls behavior. DAWN-SRW supplies precise RW-operator intervention targets, but explicit units do not guarantee semantic interpretability and

localization alone is not causal confirmation. The present evidence therefore preserves a discovery freeze followed by suppression, matched controls, and exact same-example restoration, and does not generalize the IOI result beyond the specified task, model, and intervention.

Across these lines, state dependence, conditional routing, depth sharing, attention-side experts, rank-1 RW structure, and programming interpretations all have clear precedents. The contribution here is a constructive architectural test of organizing transition-producing computation into identifiable learned operations and explicit state-conditioned control, instantiated consistently across Q/K/V/RST and evaluated for viability, conditional organization, and causal accessibility.

3. DAWN-SRW

3.1. State Evolution as the Starting Abstraction

Let

$$X_\ell = (x_{\ell,1}, \dots, x_{\ell,T}) \in \mathbb{R}^{T \times D}$$

denote the sequence representation at depth ℓ . At a coarse level, a Transformer-like block evolves this state through a relation-dependent attention update followed by a transformation of the contextualized representation:

$$Y_\ell = X_\ell + A_\ell(X_\ell), \quad (4)$$

$$X_{\ell+1} = Y_\ell + R_\ell(Y_\ell). \quad (5)$$

This decomposition only fixes the computational abstraction and notation. DAWN-SRW concerns how the transformations producing these updates are parameterized.

3.2. Explicit State-Conditioned Operator Parameterization

The architectural hypothesis is broader than any particular primitive, control geometry, or application rule. It separates identifiable learned operations from explicit state-conditioned control that determines how those operations contribute to the subsequent state transition. The evaluated v4172 architecture instantiates that hypothesis through reusable RW primitives, an applicability profile, and normalized additive composition. For a token state $x \in \mathcal{S} = \mathbb{R}^D$, let $a \in \{Q, K, V, \text{RST}\}$ index the computation type and define

$$F_a : \mathcal{S} \rightarrow \mathbb{R}^D. \quad (6)$$

The Q, K, and V outputs form the state-conditioned interfaces through which relational interaction is computed, whereas F_{RST} directly specifies a residual-state update. Each composed output is explicitly decomposed into contributions from reusable RW primitives under state-dependent control. This RW factorization, its rank-1 granularity, the

operation-address coupling, spherical geometry, and normalized additive composition are design choices of v4172. They provide one concrete test of the hypothesis rather than defining it.

3.3. Read-Write Computational Primitives

Each primitive owns learned read and write parameters. With $\text{unit}(z) = z / \max(\|z\|_2, \epsilon)$, define normalized directions $\bar{r}_{a,i} = \text{unit}(r_{a,i})$ and $\bar{w}_{a,i} = \text{unit}(w_{a,i})$. The primitive executes

$$\mathcal{O}_{a,i}(x) = \langle x, \bar{r}_{a,i} \rangle \bar{w}_{a,i} = \bar{w}_{a,i} \bar{r}_{a,i}^\top x. \quad (7)$$

This is an explicit rank-1 RW primitive. It reads a scalar from the current state and produces a directed output contribution. The read direction specifies the extracted scalar, while the write direction specifies the direction of that contribution. For Q, K, and V, these contributions compose the interfaces used by attention; for RST, they compose a direct residual-state update. Unlike a static retrieved value, an RW primitive remains an input-dependent transformation, with its output varying with the current state through the read response. The state-to-state transition is completed only when the composed output participates in the attention-mediated or residual update.

3.4. State-Conditioned Applicability

At the abstract level, control is represented by an applicability profile over reusable RW primitives. For each computation type $a \in \{Q, K, V, \text{RST}\}$, the current state induces a profile over all N_a primitives,

$$C_a(x) = (c_{a,1}(x), \dots, c_{a,N_a}(x)), \quad (8)$$

whose coordinates specify which primitives participate and to what degree, rather than assigning the state to a single discrete branch. The evaluated v4172 realization has exact zero applicability outside learned boundaries and continuously positive applicability inside them. This is not an algebraic decomposition of the representation vector; its coordinates control the contributions of the corresponding RW primitives composed below.

In v4172, normalized coordinates $q_a(x)$ and $k_{a,i}$ compare the current state with primitive i 's learned address; Section 3.6 gives their construction. A separate direct projection produces the state-dependent DirectTau boundary:

$$\tau_a(x) = -1 + 2\sigma(xu_a + \beta_a), \quad (9)$$

$$\rho_{a,i}(x) = q_a(x)^\top k_{a,i}, \quad (10)$$

$$m_{a,i}(x) = \rho_{a,i}(x) - \tau_a(x). \quad (11)$$

The evaluated `linear_angular` rule defines primitive applicability as

$$c_{a,i}(x) = \text{clip}\left(\frac{m_{a,i}(x)}{\max(1 - \tau_a(x), 10^{-4})}, 0, 1\right). \quad (12)$$

We call $c_{a,i}(x) \in [0, 1]$ the *continuous-valued applicability coefficient* for primitive i at state x . Its exact state-dependent applicability support is

$$\mathcal{I}_a(x) = \{i : \rho_{a,i}(x) > \tau_a(x)\} = \{i : c_{a,i}(x) > 0\}. \quad (13)$$

Outside this set, a primitive contributes exactly zero; inside it, $0 < c_{a,i}(x) \leq 1$. A primitive's contribution therefore enters continuously from zero at the learned boundary. This combines state-dependent sparse applicability support with a continuous-valued applicability coefficient rather than merely assigning a dense soft mixture over every primitive.

3.5. State-Conditioned Operator Composition

Applicable primitives compose into the output $F_a(x)$:

$$Z_a(x) = \left[\max\left(\sum_i c_{a,i}(x), 1\right) \right]^{p_a}, \quad (14)$$

$$s_a = \sqrt{D/L},$$

$$F_a(x) = \frac{s_a}{Z_a(x)} \sum_i c_{a,i}(x) \mathcal{O}_{a,i}(x). \quad (15)$$

Equivalently,

$$F_a(x) = A_a(x)x, \quad (16)$$

$$A_a(x) = \frac{s_a}{Z_a(x)} \sum_i c_{a,i}(x) \bar{w}_{a,i} \bar{r}_{a,i}^\top.$$

This form makes the architectural hypothesis explicit. The model does not apply one fixed dense transformation to every state. Instead, the current representation determines continuous-valued applicability coefficients, which construct a state-dependent transformation $A_a(x)$ from reusable rank-1 RW primitives. The resulting computation is both input-conditioned and explicitly decomposed into identifiable operator contributions.

“Continuous-valued” refers to the applicability coefficient, not to continuous-time dynamics or a continuous-depth model; layer evolution remains discrete. The scalar mapping from normalized margin to applicability is continuous and piecewise linear, with nondifferentiable support boundaries.

The denominator powers are $p_Q = p_K = 0.5$, $p_V = 1.0$, and $p_{\text{RST}} = 1.2$. The fixed composition scale is $s_a = \sqrt{D/L}$, with no learned pool-strength parameter.

3.6. Operation-Address Space

The operation-address space $\mathcal{A} = \mathbb{R}^R$ is where primitive applicability is evaluated, not where RW computation itself is executed. This addressing construction is one evaluated realization of the state-conditioned applicability abstraction above. The present experiments do not establish that this particular realization is necessary or preferable to alternative

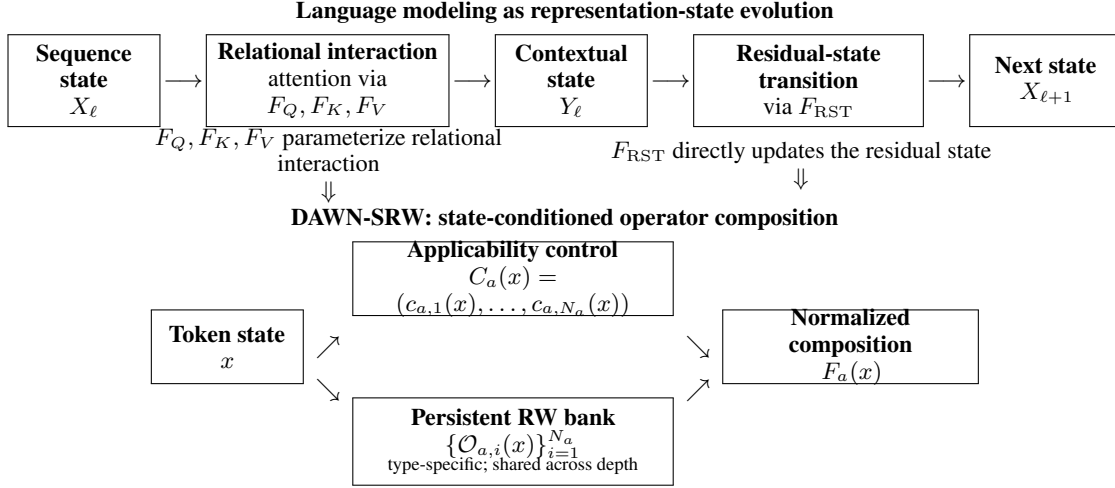


Figure 1. The two-level construction. Top: a language-model block evolves a sequence representation through relational interaction and a subsequent residual-state transition. Bottom: x separately determines the applicability profile and input-dependent responses of a persistent, type-specific RW bank shared across depth; the two meet in normalized composition $F_a(x)$. In v4172, $q_a(x)$, $\tau_a(x)$, $k_{a,i}$ implement applicability in address space; RW execution remains in the full D -dimensional representation space.

state-conditioned control mechanisms. In v4172, $R = 256$, and a normalized query is a direct projection of the current input state:

$$q_a(x) = \text{unit}(xP_a + b_a) \in \mathcal{A}. \quad (17)$$

A shared probe pair $P_r, P_w \in \mathbb{R}^{D \times R}$ maps all operator pools' functional read and write directions into address features. The operator address is

$$\begin{aligned} u_{a,i} &= \text{unit}(\bar{r}_{a,i}^\top P_r), & v_{a,i} &= \text{unit}(\bar{w}_{a,i}^\top P_w), \\ k_{a,i} &= \text{unit}(u_{a,i} \odot v_{a,i}) \in \mathcal{A}. \end{aligned} \quad (18)$$

This generalized coordinate-wise bilinear mapping makes address identity a learned function of what a primitive reads and where it writes, rather than an independent expert ID. It is the addressing realization used to compute the applicability coefficients in Equation 12, not the paper's definition of a state transition. Shared probes align QK, V, and RST addresses to one coordinate convention, while their operator rows remain in QK-, V-, and RST-specific pools.

Because both $q_a(x)$ and $k_{a,i}$ are unit-normalized, they lie on the address sphere $\mathbb{S}^{R-1}$. Let $\theta_{a,i}(x) = \arccos(q_a(x)^\top k_{a,i})$ and $\theta_a(x) = \arccos \tau_a(x)$. The applicability support is therefore

$$\mathcal{I}_a(x) = \{i : \theta_{a,i}(x) < \theta_a(x)\},$$

a spherical cap centered at the query $q_a(x)$ with state-dependent angular radius $\theta_a(x)$. Geometrically, $q_a(x)$ specifies the center of the applicability query, $\tau_a(x)$ specifies its angular support boundary, and $c_{a,i}(x)$ specifies the continuous-valued coefficient with which primitive i applies. When the numerical guard in Equation 12 is inactive,

an active primitive has

$$c_{a,i}(x) = \frac{\cos \theta_{a,i}(x) - \tau_a(x)}{1 - \tau_a(x)},$$

which rises continuously from zero at the cap boundary to one at its center. Equivalently, $c_{a,i}(x)$ is a normalized continuous depth within the cap. Equation 12 is the exact implementation, retaining the 10^{-4} denominator guard and clipping.

3.7. Attention-Mediated and Residual Transitions

For a sequence state X_ℓ , layer-specific normalization gives $\tilde{X}_\ell = \text{LN}_\ell^{(1)}(X_\ell)$. Three state-conditioned operator compositions produce the attention interfaces:

$$Q_\ell = F_Q(\tilde{X}_\ell),$$

$$K_\ell = F_K(\tilde{X}_\ell),$$

$$V_\ell = F_V(\tilde{X}_\ell),$$

$$A_\ell = \text{softmax}\left(\frac{Q_\ell K_\ell^\top}{\sqrt{d_h}} + M_{\text{causal}}\right),$$

$$\Delta X_{\text{attn},\ell} = A_\ell V_\ell W_\ell^O, \quad (19)$$

$$Y_\ell = X_\ell + \Delta X_{\text{attn},\ell}. \quad (20)$$

Q and K produce separate outputs from the same 4,600-row QK pool; V produces an output from a separate 14,500-row pool. After attention, a second layer-specific normalization $\tilde{Y}_\ell = \text{LN}_\ell^{(2)}(Y_\ell)$ is evaluated anew using the 29,300-row RST operator pool:

$$X_{\ell+1} = Y_\ell + F_{\text{RST}}(\tilde{Y}_\ell). \quad (21)$$

The block thus applies an attention-mediated state transition followed by a direct residual-state transition. RST applicability depends on the contextualized post-attention state, not the pre-attention query.

3.8. Cross-Depth Operator Reuse

The QK, V, and RST read/write pools; the two RW-to-address probes; the Q/K/V query projection; the RST query projection; and their DirectTau projections are shared across all L layers. Q and K share read/write rows but retain separate query slices, tau values, applicability coefficients, and outputs. The layer-specific parameters are two LayerNorms and the attention output projection W_ℓ^O . Positional and token embeddings and the final normalization are model-level parameters. Consequently, DAWN-SRW reuses RW primitives across depth without claiming complete block sharing.

3.9. Implemented Addressing and Compute Boundary

For computation type a , define *operator visibility* $\mathcal{V}_a(x)$ as the set of operator identities made available to the control mechanism for applicability evaluation at state x . Visibility denotes the candidate range available to that mechanism, not whether operators are observable to a human analyst. In the evaluated v4172 realization, every operator in the type-specific pool is visible:

$$\mathcal{I}_a(x) \subseteq \mathcal{V}_a(x) = \{1, \dots, N_a\},$$

with $\mathcal{I}_a(x)$ generally a strict subset.

The implementation shards operator rows across the model mesh and processes them in dense chunks. The 400M configuration uses one QK chunk, one V chunk, and four RST chunks. Every query scores the full local address pool; every operator row participates in dense read and write contractions before cross-shard reduction. Exact applicability support can be sparse while operator visibility remains global and RW execution remains dense; it is not an indexed sparse kernel.

The evaluated architecture keeps RW execution in the full residual dimension D ; the R -dimensional operation-address space is used only for applicability evaluation. An executor could omit exactly zero RW terms without changing the mathematical output, but such support-restricted RW execution is not implemented or timed here.

4. Experimental Setup

4.1. Models and training

Table 1 summarizes the evaluated checkpoints. The primary DAWN-SRW model has 393,800,708 measured parameters, operation-address dimension $R = 256$, and 48,400

operator rows across QK, V, and RST. It was trained for 39,999,504,384 packed C4 tokens with sequence length 512 and global training batch 1,024. Both dropout rates are zero.

4.2. Matched 400M dense control

The dense control has 393,896,960 measured parameters, width 1,280, feed-forward width 5,120, 18 layers, and 20 heads. It uses the same tokenizer, vocabulary, sequence length, training-token count, seed, batch size, mesh, learning rate, warmup ratio, weight decay, and zero dropout as the primary model. The parameter difference is 96,252, or 0.0244% of the dense model. This comparison tests performance retention under matched scale and training exposure; it is not a formal equivalence test.

4.3. Evaluation protocols

The C4 protocol evaluates 610 complete global batches from the fixed 10M-token capped packed-validation view: 19,520 sequences and 9,974,720 valid next-token targets. Cross-entropy, perplexity, and next-token accuracy use float32 scoring. The zero-shot protocol uses `lm-eval` 0.4.2 with no few-shot examples and no task limit on LAMBADA, HellaSwag, PIQA, ARC-Easy, ARC-Challenge, and WinoGrande. The primary metric is raw accuracy for LAMBADA and WinoGrande and normalized accuracy for the other four tasks. The mean is an unweighted mean of these six metrics.

4.4. Mechanistic localization and frozen intervention protocol

We define a site by layer, computation type, token position, and global operator identity. IOI discovery uses 119 independent paired-correct units. Sites are ranked by absolute discovery mean contribution importance. We choose the smallest audited prefix that reaches at least 0.70 cumulative absolute importance and 0.95 split top- k overlap. The resulting 4,096-site Q/K/V/RST/layer composition and intervention specification are frozen before validation.

The primary causal score is the summed-log-probability margin between the correct indirect object and the source name. Circuit-wide suppression zeros only the selected operators' execution numerator at target positions while retaining the full production denominator. Four control families each contain 100 frozen-size circuits that exclude circuit sites: equal-count random, same-layer random, activation-matched, and route-frequency-matched across Q/K/V/RST. Exact restoration reinserts the same-example selected numerator after suppression, again under the production denominator. Validation uses 123 independent paired-correct units; the held-out test accessor is opened only after the validation result and unchanged specification are finalized. We report 95% bootstrap confidence intervals and two-sided

Table 1. Model and training specification. The 1.3B row is descriptive and does not form a controlled scaling curve with the 400M rows.

Model	Parameters	Tokens	D	Layers	Heads	QK / V / RST rows	Checkpoint
DAWN-SRW 400M	393,800,708	39,999,504,384	2,304	18	36	4,600 / 14,500 / 29,300	76,293
Dense Transformer 400M	393,896,960	39,999,504,384	1,280	18	20	—	76,293
DAWN-SRW 1.3B	1,272,038,404	19,999,981,568	4,096	24	32	8,534 / 26,904 / 54,352	87,193

Table 2. Packed C4 validation on the identical 9,974,720 target tokens.

Model	CE	PPL	Accuracy
DAWN-SRW 400M	2.940213	18.919873	0.424912
Dense 400M	2.935894	18.838336	0.427100
DAWN-SRW 1.3B	2.835879	17.045382	0.437099

paired permutation tests where specified.

5. Results

5.1. Language-model performance retention at 400M

Table 2 reports the locked C4 evaluation. DAWN-SRW differs from the dense control by $+0.004318918$ cross-entropy, $+0.081537182$ perplexity, and -0.002188432 accuracy (-0.2188 percentage point). These small measured gaps support performance retention, not superiority or formal equivalence.

Table 3 gives task-level zero-shot results with standard errors. The matched 400M means are 0.462872 and 0.459652, a difference of $+0.003220495$ ($+0.3220$ percentage point). Task differences vary in sign, so the mean does not establish benchmark superiority. LAMBADA perplexities are 24.0774 ± 0.8005 for DAWN-SRW 400M, 24.0042 ± 0.8073 for dense 400M, and 17.6358 ± 0.5390 for DAWN-SRW 1.3B.

5.2. Descriptive larger-scale evaluation at 1.3B

The 1.272B checkpoint reaches 2.835879 C4 cross-entropy, 17.045382 perplexity, 0.437099 token accuracy, and a 0.485844 six-task mean. It was trained for 19.999B tokens, about half the 400M training-token count, and no completed matched dense 1.3B result is admitted to this study. These results therefore describe a larger trained DAWN-SRW model; they are neither a controlled scaling curve nor a dense-model comparison.

5.3. Input-conditioned RW applicability support

At all 9,994,240 packed-C4 positions, Q, K, V, and RST each have strict-subset applicability support at positive selection margin (Table 4). Mean exact applicability support is 487.896 Q operators, 386.722 K operators, 812.388 in the shared-pool Q/K union, 1,649.994 V operators, and

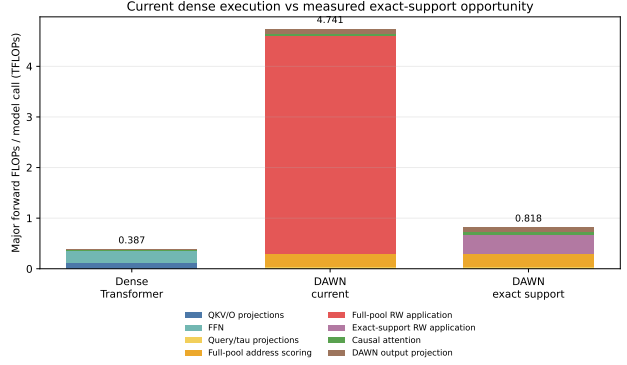


Figure 2. Static backbone accounting at batch 1 and sequence length 512. Address scoring remains dense in the estimate restricted to exact applicability support. Current execution also applies all full-dimensional RW rows. No production latency claim is made.

2,017.703 RST operators. This is exact input-conditioned applicability support under Equation 12.

Applicability support does not make the present kernel efficient. Static forward accounting at batch 1 and sequence length 512 is 4.741260116 TFLOPs for current dense DAWN-SRW, 0.818230250 TFLOPs for dense addressing plus RW application restricted to exact applicability support, and 0.386547057 TFLOPs for the matched dense Transformer. The corresponding opportunity is 17.2577% of current DAWN-SRW accounting but still 2.1168 times the dense Transformer (Figure 2). It is a mathematical conditional-RW opportunity, not measured latency, energy, or hardware speedup.

5.4. IOI circuit localization

Discovery produces split-half rank stability 0.9139788 over 456,697 ranked sites. The smallest prefix passing both frozen discovery gates contains 4,096 sites and 0.719249 cumulative absolute importance. It spans all 18 layers and contains 1,744 Q sites, 2,350 K sites, two V sites, and no RST sites. Split top- k overlap is 0.974854 (Jaccard 0.950941). These statistics identify a Q/K-centered candidate circuit; localization alone is not a causal claim.

5.5. Validation and held-out causal intervention

Table 5 shows confirmatory results. On validation, suppression moves the mean margin from 3.465456 to -4.399331 ,

Table 3. Zero-shot accuracy. Parentheses contain standard errors. HS, AE, and AC denote HellaSwag, ARC-Easy, and ARC-Challenge.

Model	LAMBADA	HS	PIQA	AE	AC	WinoGrande	Mean
DAWN-SRW 400M	0.352028 (0.006654)	0.458375 (0.004972)	0.712187 (0.010563)	0.484007 (0.010255)	0.256826 (0.012767)	0.513812 (0.014047)	0.462872
Dense 400M	0.353192 (0.006659)	0.457080 (0.004971)	0.698041 (0.010712)	0.475589 (0.010248)	0.249147 (0.012639)	0.524862 (0.014035)	0.459652
DAWN-SRW 1.3B	0.395886 (0.006813)	0.510157 (0.004989)	0.723069 (0.010440)	0.509680 (0.010258)	0.268771 (0.012955)	0.507498 (0.014051)	0.485844

Table 4. Exact applicability support at selection margin > 0 . Q/K union counts an operator active for Q or K at the same layer and token.

Computation	Pool	Mean active	Pool fraction
Q	4,600	487.896	10.6064%
K	4,600	386.722	8.4070%
Q/K union	4,600	812.388	17.6606%
V	14,500	1,649.994	11.3793%
RST	29,300	2,017.703	6.8864%

a drop of 7.864787 (95% CI [7.490080, 8.241263]), and accuracy falls from 1.000000 to 0.008130. The frozen effect exceeds every replicate in all four 100-circuit control families. The most competitive control is activation-matched, with mean drop 5.400490; the frozen-minus-control difference is 2.464297 (95% CI [2.076200, 2.832697], paired permutation $p = 0.00049975$). Exact restoration returns accuracy to 0.991870 and recovers 0.999661 of the suppression effect (95% CI [0.998655, 1.000584]).

The fully held-out test reuses the unchanged site identity, scoring rule, intervention, controls, seeds, and thresholds. Suppression moves margin from 3.254720 to -4.155613 , a drop of 7.410332 (95% CI [7.016485, 7.818472]), while accuracy falls to 0.008772. Again the frozen effect exceeds every replicate in all four control families. The activation-matched mean drop is 5.472579; the frozen-minus-control difference is 1.937754 (95% CI [1.564818, 2.319706], $p = 0.00049975$). Restoration recovers 0.998406 of the effect (95% CI [0.997324, 0.999385]) and restores accuracy to 1.000000.

Together, suppression, matched-control separation, and exact restoration support the claim that this discovery-frozen Q/K-centered operator circuit is causally necessary for the evaluated IOI behavior on independent validation and fully held-out inputs. “Necessary” is intervention-specific: it does not mean the sites form a unique or human-readable algorithm.

5.6. Negative and fail-closed results

The same workflow produces informative negative boundaries (Table 6). RAVEL passes the behavioral gate but fails variable-localization stability, so no circuit is frozen and no confirmatory split is opened. ARC discovery freezes a stable 4,096-site Q/K circuit, but validation suppression has the wrong sign and fails control separation; restoration passes,

and test remains closed. For 1.3B IOI, behavioral eligibility and rank stability pass, but the largest allowed 4,096-site prefix reaches only 0.635144 cumulative importance, below the preregistered 0.70 threshold. No circuit is frozen and no validation or test is opened.

6. Discussion

6.1. What the results establish

The constructive test supports three ordered claims. First, the matched 400M comparison establishes viability under the measured protocols. Second, exact input-conditioned applicability support establishes that control is state-dependent in the evaluated model, even though visibility and execution remain dense. Third, the discovery-frozen IOI result establishes causal accessibility for the specified circuit: suppressing its operator contributions changes behavior, matched controls do not reproduce that effect, and exact same-example restoration recovers it on independent validation and held-out data. RAVEL, ARC, and 1.3B IOI are fail-closed boundaries showing that successful localization is not automatic across tasks or scales.

These results do not establish that language reasoning is a discrete program, that the factorization is intrinsic or optimal, or that identifiable operations have human-readable semantics. They show that one explicit factorization can support language modeling, exhibit state-conditioned organization, and provide causal intervention targets under the evaluated conditions.

6.2. From Architectural Hypothesis to an Open Design Space

Returning to the broad question posed at the outset, the architectural hypothesis separates identifiable learned operations from explicit state-conditioned control that determines how those operations contribute to the subsequent state transition. “Identifiable” is an architectural property: an operation is referable as a stable object that can be measured and intervened on. It does not prescribe the scope or duration of that identity and is not a claim of semantic purity. Operations may remain distributed or polysemantic, and a behavior may depend on many operation sites.

v4172 occupies one point in this design space. It represents operations as rank-1 RW primitives; keeps type-specific QK, V, and RST banks persistent across depth; couples operation

Table 5. Frozen IOI circuit results. Accuracy arrows show intact to suppressed. Recovery is the fraction of the suppression effect recovered by exact same-example numerator restoration.

Split	Units	Intact margin	Suppressed margin	Drop (95% CI)	Accuracy	Recovery / restored acc.
Validation	123	3.465456	-4.399331	7.864787 [7.490080, 8.241263]	1.000000 → 0.008130	0.999661 / 0.991870
Held-out test	114	3.254720	-4.155613	7.410332 [7.016485, 7.818472]	1.000000 → 0.008772	0.998406 / 1.000000

Table 6. Fail-closed mechanistic results. Later phases were not opened after a failed preregistered gate.

Study	Passed	Failed boundary	Decision
RAVEL 400M	Behavioral eligibility	Variable stability: 0.628762 / 0.676243 / 0.660622	No frozen circuit; no validation, transfer, or test
ARC 400M	Stable discovery; 4,096 Q/K sites; restoration	Suppression effect wrong sign; control separation failed	Reject causal replication; held-out test closed
IOI 1.3B	Behavioral eligibility; rank stability 0.907188	4,096-site cumulative importance 0.635144 < 0.70	No frozen circuit; validation and test closed

identity and control through read/write-derived addresses in a learned spherical geometry; uses state-derived queries and thresholds to produce exact support with continuous-valued coefficients; and applies normalized additive composition across Q/K/V/RST. These choices make the objects and their contributions explicit, but none is required by the higher-level hypothesis. The programming analogy and the view of depth as state evolution likewise organize the investigation; they do not supply a formal semantics for the model.

The hypothesis therefore opens several independent design questions. Operation representation and granularity may differ from rank-1 RW units. Control may couple or separate functional parameters and identity, and operation organization may be flat or hierarchical and persistent over different scopes. Application semantics need not be normalized addition: alternative systems could order, sequence, or otherwise combine controlled operations. Finally, operator visibility may be restricted before applicability is evaluated, while physical execution may separately exploit the resulting conditional structure. The present analysis interface supports tracing, suppression, and restoration of explicit contributions, but reliable semantic editing, transfer, and human-readable programming remain open goals rather than established capabilities.

6.3. From Global Operator Visibility to Candidate-Restricted Execution

The present results establish state-conditioned applicability support: for a given representation state, many operators have exactly zero applicability. They do not establish candidate-restricted operator visibility. In v4172, $\mathcal{V}_a(x)$ remains the full type-specific pool, so applicability is evaluated before zero-contribution operators are excluded. A stronger execution path would first identify a substantially smaller visible candidate set, then evaluate applicability and compose within that set.

This is the most immediate execution question raised by

v4172. Identification or retrieval would first restrict operator visibility; applicability would still determine applicability support within the visible set; accelerator-aware dispatch would then need to avoid evaluating inactive candidates. The central empirical question is how far operator visibility and executed computation can be reduced while preserving language-model quality and expressivity.

6.4. Limitations

The present results do not establish that transition-producing computation has been factorized into an efficient conditional computational structure. Although each state has strict-subset applicability support, applicability is still evaluated with global operator visibility over large type-specific pools. Performance retention may therefore depend partly on the representational capacity provided by a large distributed operator bank and full-pool condition evaluation, rather than on a compact or intrinsically efficient decomposition of language-model computation.

Relatedly, the learned spherical address geometry is not shown to reflect an intrinsic or compressible organization of language operations. It may instead serve as a high-capacity substrate in which useful operators are distributed broadly enough for global state-conditioned access to recover the required transformations. The present study therefore does not establish how much operator visibility is required to preserve language-model expressivity, or whether candidate evaluation and executed computation can be restricted substantially further.

The current implementation evaluates full address and operator pools densely. Exact applicability support does not imply current sparse execution, and even the accounting restricted to exact applicability support remains 2.1168 times the matched dense Transformer’s backbone accounting. No production latency, energy, or hardware advantage is measured.

The central causal evidence is confined to IOI and does not automatically generalize to every task. RAVEL, ARC, and the selected 1.3B IOI replication provide valid negative boundaries. The 400M and 1.3B checkpoints have different parameter and training-token budgets, so they do not define a controlled scaling curve. Operator-level causal units need not map one-to-one to human-readable concepts. Finally, the present work does not establish language reasoning as universally discrete operator programs.

7. Conclusion

This paper asks whether learned language-model computation can be organized as an explicit state-conditioned operational system. It tests a narrower architectural hypothesis: transition-producing computation can be factorized into identifiable learned operations and explicit state-conditioned control that determines how those operations contribute to the subsequent state transition.

v4172 is one concrete realization, using persistent rank-1 RW banks, read/write-derived spherical addresses, exact applicability support with continuous-valued coefficients, and normalized additive composition across Q/K/V/RST. The empirical results provide a bounded constructive test: measured 400M viability, descriptive 1.3B evidence, input-conditioned applicability support, and causal accessibility through a discovery-frozen IOI circuit with suppression, matched controls, and exact restoration. Fail-closed results mark where the same localization program did not succeed.

These findings do not establish an intrinsic factorization, human-readable operator semantics, general editability, benchmark superiority, or current hardware efficiency. They instead motivate an open research program spanning operation representation and granularity, control and identity, application semantics and persistence, and the separation of operator visibility from physical execution. The immediate question is how much global visibility and executed computation can be reduced while preserving language-model quality and expressivity.

A. Exact Addressing and Composition Equations

This appendix restates the production computation as a compact reference. Raw read and write parameters are normalized in float32 before being cast for dense contractions:

$$\bar{r}_i = \frac{r_i}{\max(\|r_i\|_2, \epsilon)}, \quad \bar{w}_i = \frac{w_i}{\max(\|w_i\|_2, \epsilon)}.$$

One pair of model-level probes is shared across QK, V, and RST:

$$u_i = \text{unit}(\bar{r}_i^\top P_r), \quad v_i = \text{unit}(\bar{w}_i^\top P_w), \\ k_i = \text{unit}(u_i \odot v_i).$$

The query and threshold use separate direct projections:

$$\hat{q}_a = \text{unit}(xP_a + b_a), \\ \tau_a = -1 + 2\sigma(xu_a + \beta_a), \quad \rho_i = \hat{q}_a^\top k_i.$$

For linear angular composition:

$$c_i = \text{clip}\left(\frac{\rho_i - \tau_a}{\max(1 - \tau_a, 10^{-4})}, 0, 1\right)$$

and padding rows are masked to zero. This is the single applicability coefficient used in RW composition. Let $M_a = \sum_i c_i$. The production composed output before causal attention or residual addition is

$$F_a(x) = \sqrt{\frac{D}{L}} \frac{\sum_i c_i \langle x, \bar{r}_i \rangle \bar{w}_i}{\max(M_a, 1)^{p_a}},$$

with $p_Q = p_K = 0.5$, $p_V = 1.0$, and $p_{RST} = 1.2$. Operator rows are sharded over the model axis; local numerators and admission mass are reduced across shards. Chunking changes the contraction schedule, not the equation.

B. Complete Model and Training Settings

The 400M DAWN-SRW checkpoint is the primary performance and mechanistic model. The dense 400M checkpoint is only a matched performance control. The 1.3B DAWN-SRW checkpoint is used for descriptive performance and one fail-closed IOI localization replication. The still-incomplete dense 1.3B run is excluded.

C. Benchmark Metrics and Protocol Details

C4. All reported C4 rows use the same packed validation artifact and `packed_c4_validation_v1` protocol. The capped view contains 19,531 sequences. Exactly 19,520 sequences form 610 complete global evaluation batches, producing 9,974,720 valid next-token targets; 11 incomplete-tail sequences are excluded before scoring.

Zero-shot. The protocol is `mamba-table3_zero_shot_v1`, using `stock lm-eval==0.4.2`, global batch 32, maximum length 512, length buckets 64/128/256/512, and zero few-shot examples. Evaluated split sizes are 5,153 for LAMBADA, 10,042 for HellaSwag, 1,838 for PIQA, 2,376 for ARC-Easy, 1,172 for ARC-Challenge, and 1,267 for WinoGrande. Task, dataset, tokenizer, and software contracts match exactly between the two 400M models.

D. IOI Localization and Intervention Details

D.1. Frozen discovery

Discovery admits 119 paired-correct independent units and ranks 456,697 sites. Split-half rank stability is 0.913978765.

Table 7. Detailed settings from the registered training configurations and restored checkpoint manifests.

Setting	DAWN-SRW 400M	Dense 400M	DAWN-SRW 1.3B
Vocabulary	30,522	30,522	30,522
Residual width	2,304	1,280	4,096
Layers	18	18	24
Heads / head width	36 / 64	20 / 64	32 / 128
Maximum sequence length	512	512	512
Feed-forward width	—	5,120	—
Operation-address width	256	—	256
QK / V / RST pool	4,600 / 14,500 / 29,300	—	8,534 / 26,904 / 54,352
QK / V / RST chunks	1 / 1 / 4	—	1 / 2 / 8
QK / V / RST denominator power	0.5 / 1.0 / 1.2	—	0.5 / 1.0 / 1.2
Initial active targets	0.05 / 0.03 / 0.01	—	0.05 / 0.03 / 0.01
Training batch	1,024	1,024	448
Mesh (data $\times$ model)	16×2	16×2	16×2
Learning rate	3×10^{-4}	3×10^{-4}	2×10^{-4}
Warmup ratio	0.08	0.08	0.06
Weight decay	0.1	0.1	0.1
Tau learning-rate multiplier	0.001	—	0.001
Dropout / <code>router_dropout</code>	0 / 0	0 / —	0 / 0
Gradient checkpointing	yes	yes	yes
Seed	1	1	1

The frozen selection rule is the smallest prefix satisfying cumulative absolute importance at least 0.70 and split top- k overlap at least 0.95. The selected $k = 4096$ prefix has:

- cumulative absolute importance 0.719248713;
- 3,993 shared split sites, overlap 0.974853516, and Jac-card 0.950940700;
- Q/K/V/RST counts = 1744/2350/2/0;
- coverage of every layer from 0 through 17.

Validation and test do not recompute or alter this selection.

D.2. Suppression, controls, and restoration

For a correct candidate c and source candidate s , the primary score is

$$\mathcal{M} = \sum_{t \in c} \log p(t) - \sum_{t \in s} \log p(t).$$

Suppression zeros the frozen sites’ execution numerator at specified token positions. It leaves admission mass and the production denominator unchanged. Controls use the same 4,096-site count and exclude all circuit sites. Activation matching conditions on layer, computation type, and the discovery mean pre-scale operator-output-norm quantile. Restoration executes a distinct intervention: after suppression, it adds back the exact selected same-example numerator under the unchanged denominator.

On validation, restoration returns the margin to 3.462791, leaving residual drop 0.002665. On held-out test, it returns the margin to 3.242905, leaving residual drop 0.011815. The respective recovery fractions are 0.999661155 (95% CI [0.998654544, 1.000584125]) and 0.998405645 (95% CI [0.997323737, 0.999385266]).

E. Negative Mechanistic Results

RAVEL. Discovery uses 142 independent units. Rank stability is 0.628762 for Continent, 0.676243 for Country, and 0.660622 for Language, below the preregistered requirement. No prefix passes the complete discovery rule, no variable circuit is frozen, and validation, transfer, and test are not opened.

ARC. Discovery stability is 0.89386 and a 4,096-site Q/K-only circuit is frozen. On 51 validation units, intact margin is 1.735534 and suppressed margin is 2.819139, giving a frozen margin drop of -1.083605 (95% CI $[-1.472180, -0.690348]$). The effect has the wrong sign and fails control separation. Restoration recovery is 1.004083 (95% CI [0.994587, 1.013652]), but restoration alone does not pass the conjunction. Held-out test remains closed.

1.3B IOI. The scale replication admits 125 independent units and passes rank stability at 0.907187749. The largest allowed 4,096-site prefix reaches only 0.635144 cumulative absolute importance, below the 0.70 gate. No circuit is frozen; validation and test remain closed.

F. Compute Accounting

The applicability-support analysis evaluates 9,994,240 packed-C4 positions and observes address-stage margins before RW application. It stores aggregate statistics, not per-token operator identities or vectors. Static accounting uses batch 1, sequence length 512, and 1 MAC = 2 FLOPs. It counts major matmul/einsum operations and omits softmax, normalization, gate elementwise arithmetic, embedding lookup, residual additions, and communication collectives.

The estimate restricted to exact applicability support is mathematically output-equivalent if the executor can skip all and

Table 8. Control effects and frozen-minus-control separation. Each control family contains 100 replicates. All paired permutation p -values are 0.0004997501 with 2,000 permutations.

Control	Validation drop	Frozen minus control (95% CI)	Test drop	Frozen minus control (95% CI)
Equal-count	0.070886	7.793902 [7.437851, 8.161572]	0.080276	7.330057 [6.922408, 7.716879]
Same-layer	0.039389	7.825399 [7.440629, 8.180288]	0.041073	7.369260 [6.977447, 7.754539]
Activation-matched	5.400490	2.464297 [2.076200, 2.832697]	5.472579	1.937754 [1.564818, 2.319706]
Route-frequency	0.219162	7.645626 [7.289512, 8.031625]	0.242665	7.167668 [6.779836, 7.545537]

Table 9. Backbone static accounting. “Exact applicability support” retains full dense address scoring and replaces only full-pool RW application with the measured positive-margin applicability support.

Component	Current dense DAWN-SRW	Dense address + exact RW	Dense Transformer
Address scoring + query/tau projections	0.293742	0.293742	–
RW application	4.306187	0.383157	–
Causal attention	0.043487	0.043487	0.024159
Dense output / QKV / FFN maps	0.097845	0.097845	0.362388
Total TFLOPs	4.741260	0.818230	0.386547

only zero-weight RW terms. It does not include indexed address retrieval and does not demonstrate kernel parity. Instrumented profiling time is excluded from latency claims because histogram collection is part of the measured path.

G. Reproducibility Identities

The machine-readable checkpoint registry is `configs/paper_checkpoint_registry.yaml`. Exact evaluation commits, software versions, data hashes, result roots, and complete artifact hashes are recorded in the repository experiment ledger. Historical experiment branches are evidence identities only; this source is maintained on the canonical main branch.

Acknowledgments

The author thanks Anthropic Claude, Claude Code, OpenAI ChatGPT, and OpenAI Codex for assistance throughout this research. This work was developed with extensive use of AI systems for literature exploration, conceptual exploration, cross-domain framing, abstraction and terminology refinement, implementation support, debugging, and manuscript preparation. In particular, these systems were used as interactive assistants for refining the learned state-transition formulation and the exposition of DAWN-SRW. The research direction, architectural decisions, experimental design, experiments, analyses, interpretations, and final claims remain the responsibility of the author.

This research was supported by the Google TPU Research Cloud (TRC) program, which provided computational resources for training and analysis.

References

- Berges, V.-P., Oğuz, B., Haziza, D., Yih, W.-t., Zettlemoyer, L., and Ghosh, G. Memory layers at scale. *arXiv preprint arXiv:2412.09764*, 2024. doi: 10.48550/arXiv.2412.09764.
- Chen, Y., Gu, N., Shang, J., Zhang, Z., Feng, Y., Sheng, J., Liu, T., Wang, S., Sun, Y., Wu, H., and Wang, H. Mixture of universal experts: Scaling virtual width via depth-width transformation. *arXiv preprint arXiv:2603.04971*, 2026. doi: 10.48550/arXiv.2603.04971.
- Cloutier, J. R. State-dependent riccati equation techniques: An overview. In *Proceedings of the 1997 American Control Conference*, volume 2, pp. 932–936, 1997. doi: 10.1109/ACC.1997.609663.
- Csordás, R., Irie, K., Schmidhuber, J., Potts, C., and Manning, C. D. MoEUT: Mixture-of-experts universal transformers. In *Advances in Neural Information Processing Systems*, volume 37, 2024a. doi: 10.52202/079017-0897.
- Csordás, R., Piekos, P., Irie, K., and Schmidhuber, J. SwitchHead: Accelerating transformers with mixture-of-experts attention. In *Advances in Neural Information Processing Systems*, volume 37, 2024b. doi: 10.52202/079017-2368.
- Dehghani, M., Gouws, S., Vinyals, O., Uszkoreit, J., and Kaiser, Ł. Universal transformers. In *International Conference on Learning Representations*, 2019.
- Fedus, W., Zoph, B., and Shazeer, N. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39, 2022.
- Foerster, J. N., Gilmer, J., Sohl-Dickstein, J., Chorowski, J., and Sussillo, D. Input switched affine networks: An RNN architecture designed for interpretability. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pp. 1136–1145, 2017.
- Geva, M., Schuster, R., Berant, J., and Levy, O. Transformer feed-forward layers are key-value memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 5484–5495, 2021. doi: 10.18653/v1/2021.emnlp-main.446.
- Ha, D., Dai, A., and Le, Q. V. Hypernetworks. In *International Conference on Learning Representations*, 2017.
- He, X. O. Mixture of a million experts. *arXiv preprint arXiv:2407.04153*, 2024. doi: 10.48550/arXiv.2407.04153.
- Huang, J., Wu, Z., Potts, C., Geva, M., and Geiger, A. RAVEL: Evaluating interpretability methods on disentangling language model representations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 8669–8687, 2024. doi: 10.18653/v1/2024.acl-long.470.
- Huang, M., Shi, H., Zheng, C., Wu, Y., Chen, G., Yu, X., Yin,

Table 10. Primary immutable identities. Branch names are omitted because they move; exact checkpoint steps, hashes, and artifact locations are the binding identities.

Evidence	Exact identity	Repository source
DAWN-SRW 400M	model spatial-r1-v4.1.7.2; step 76,293; checkpoint identity a7ce8afc...f91c17; config hash 08733ae4...5039b	configs/train_config_v4172_400M_c4_40B_v4_64_ver1_den_qk0p5_vlp0_rst1p2.yaml
Dense 400M	model baseline-JAX; step 76,293; config hash c3ce5d72...82fc25	configs/train_config_baseline_tpu_400M_c4_40B_v4_64_tp2.yaml
DAWN-SRW 1.3B	model spatial-r1-v4.1.7.2; step 87,193; checkpoint identity 2e38b41f...bf5f9; config hash 1e877263...4a5222	configs/train_config_v4172_1p3B_c4_20B_v4_64_ver1_den_qk0p5_vlp0_rst1p2.yaml
IOI circuit	parsed frozen-spec hash 6ec26618...418c; site hash 573cfff8...0200; selected-row hash 1439f6f3...fe0c1	configs/paper_ioi_frozen_circuit_v4172_400m.yaml
Compute support	support hash b1eac5f5...bdc2; accounting hash d7ad6115...984e	assets/paper/v4172_400m_compute_support_step76293/

- Y., and Cheng, H. UniPool: A globally shared expert pool for mixture-of-experts. *arXiv preprint arXiv:2605.06665*, 2026. doi: 10.48550/arXiv.2605.06665.
- Kramár, J., Lieberum, T., Shah, R., and Nanda, N. AtP*: An efficient and scalable method for localizing LLM behaviour to components. *arXiv preprint arXiv:2403.00745*, 2024. doi: 10.48550/arXiv.2403.00745.
- Lample, G., Sablayrolles, A., Ranzato, M., Denoyer, L., and Jégou, H. Large memory layers with product keys. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Lu, H., Zhao, C., Xue, M., Yao, L., Moore, K., and Gong, D. Little by little: Continual learning via incremental mixture of rank-1 associative memory experts. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. arXiv:2506.21035.
- Neelakantan, A., Le, Q. V., and Sutskever, I. Neural programmer: Inducing latent programs with gradient descent. In *International Conference on Learning Representations*, 2016.
- Plotkin, G. D. A structural approach to operational semantics. Technical Report DAIMI FN-19, Computer Science Department, Aarhus University, Aarhus, Denmark, September 1981.
- Reed, S. and de Freitas, N. Neural programmer-interpreters. In *International Conference on Learning Representations*, 2016.
- Schlag, I., Irie, K., and Schmidhuber, J. Linear transformers are secretly fast weight programmers. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pp. 9355–9366, 2021.
- Shazeer, N., Mirhoseini, A., Maziarz, K., Davis, A., Le, Q., Hinton, G., and Dean, J. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *ICLR*, 2017.
- Shi, J., Peng, Z., Zhu, Y., Wu, Y., Liu, G., and Luo, Y. OmniMoE: An efficient MoE by orchestrating atomic experts at scale. *arXiv preprint arXiv:2602.05711*, 2026. doi: 10.48550/arXiv.2602.05711.
- Sutskever, I., Martens, J., and Hinton, G. E. Generating text with recurrent neural networks. In *Proceedings of the 28th International Conference on Machine Learning*, pp. 1017–1024, 2011.
- Wang, K., Variengien, A., Conmy, A., Shlegeris, B., and Steinhardt, J. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small. In *International Conference on Learning Representations*, 2023.
- Yang, Y., Wang, C., and Li, J. UMoE: Unifying attention and FFN with shared experts. In *Advances in Neural Information Processing Systems*, volume 38, 2025.